Data Ingestion from the RDS to HDFS using Sqoop

Sqoop Import command used for importing table from RDS to HDFS:

```
sqoop import \
--connect jdbc:mysql://upgradtest1.cyaielc9bmnf.us-east-1.rds.amazonaws.com/testdatabase \
--table SRC_ATM_TRANS \
--username student --password STUDENT123 \
--target-dir /user/root/SRC_ATM_TRANS \
-m 1
```

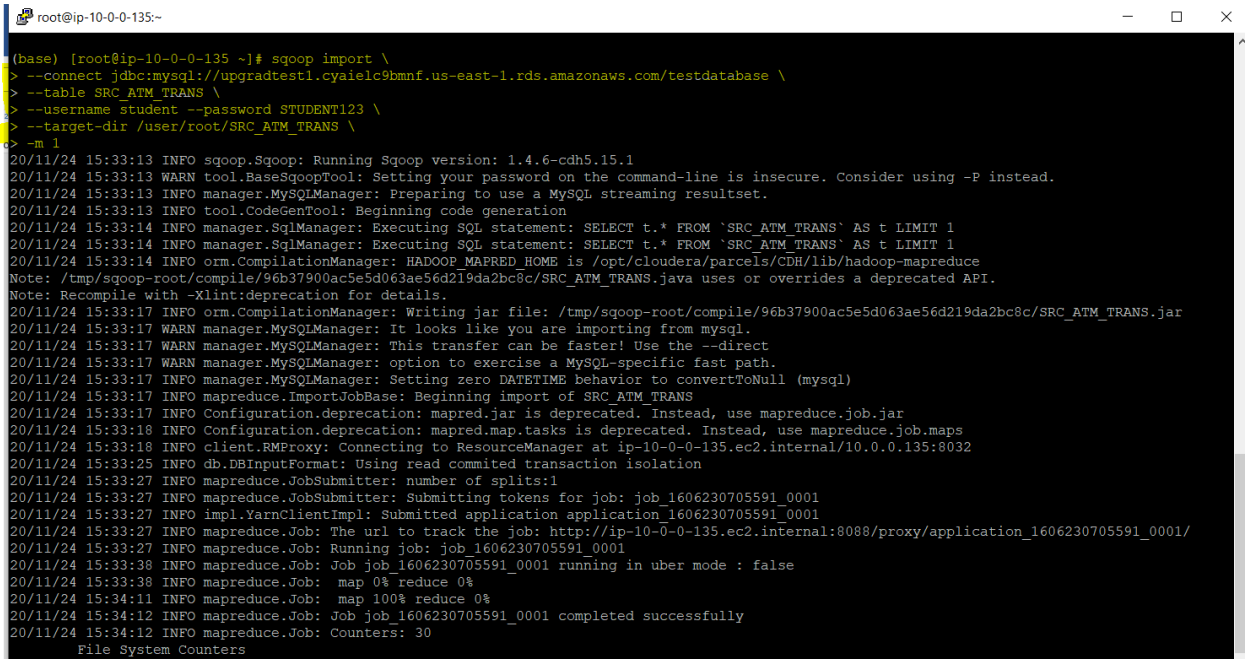

Command used to see the list of imported data in HDFS:

Hadoop fs -ls

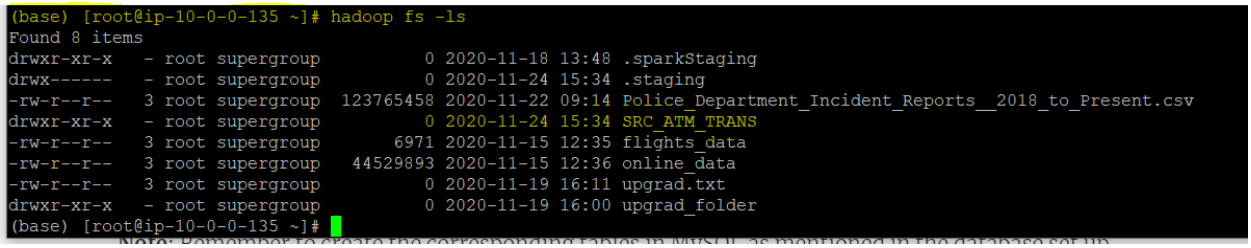

Hadoop fs -ls /user/root/SRC_ATM_TRANS

```
root@ip-10-0-0-135:~
login as: ec2-user
Authenticating with public key "imported-openssh-key"
Last login: Wed Nov 25 15:52:57 2020 from 49.35.120.124
[ec2-user@ip-10-0-0-135 ~]$ sudo -i
(base) [root@ip-10-0-0-135 ~]# hadoop fs -ls
Found 8 items
drwxr-xr-x   - root supergroup          0 2020-11-18 13:48 .sparkStaging
drwx------   - root supergroup          0 2020-11-24 15:34 .staging
-rw-r--r--   3 root supergroup  123765458 2020-11-22 09:14 Police_Department_Inc
ident_Reports__2018_to_Present.csv
drwxr-xr-x   - root supergroup          0 2020-11-24 15:34 SRC_ATM_TRANS
-rw-r--r--   3 root supergroup       6971 2020-11-15 12:35 flights_data
-rw-r--r--   3 root supergroup   44529893 2020-11-15 12:36 online_data
-rw-r--r--   3 root supergroup          0 2020-11-19 16:11 upgrad.txt
drwxr-xr-x   - root supergroup          0 2020-11-19 16:00 upgrad_folder
(base) [root@ip-10-0-0-135 ~]# hadoop fs -ls /user/root/SRC_ATM_TRANS
Found 2 items
-rw-r--r--   3 root supergroup          0 2020-11-24 15:34 /user/root/SRC_ATM_TR
ANS/_SUCCESS
-rw-r--r--   3 root supergroup  522660639 2020-11-24 15:34 /user/root/SRC_ATM_TR
ANS/part-m-00000
(base) [root@ip-10-0-0-135 ~]#
```

Screenshot of the imported data:

```
root@ip-10-0-0-135:~
20/11/24 15:33:27 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1606230705591_0001
20/11/24 15:33:27 INFO impl.YarnClientImpl: Submitted application application_1606230705591_0001
20/11/24 15:33:27 INFO mapreduce.Job: The url to track the job: http://ip-10-0-0-135.ec2.internal:8088/proxy/application_1
20/11/24 15:33:27 INFO mapreduce.Job: Running job: job_1606230705591_0001
20/11/24 15:33:38 INFO mapreduce.Job: Job job_1606230705591_0001 running in uber mode : false
20/11/24 15:33:38 INFO mapreduce.Job:  map 0% reduce 0%
20/11/24 15:34:11 INFO mapreduce.Job:  map 100% reduce 0%
20/11/24 15:34:12 INFO mapreduce.Job: Job job_1606230705591_0001 completed successfully
20/11/24 15:34:12 INFO mapreduce.Job: Counters: 30
        File System Counters
                FILE: Number of bytes read=0
                FILE: Number of bytes written=176683
                FILE: Number of read operations=0
                FILE: Number of large read operations=0
                FILE: Number of write operations=0
                HDFS: Number of bytes read=87
                HDFS: Number of bytes written=522660639
                HDFS: Number of read operations=4
                HDFS: Number of large read operations=0
                HDFS: Number of write operations=2
        Job Counters
                Launched map tasks=1
                Other local map tasks=1
                Total time spent by all maps in occupied slots (ms)=30934
                Total time spent by all reduces in occupied slots (ms)=0
                Total time spent by all map tasks (ms)=30934
                Total vcore-milliseconds taken by all map tasks=30934
                Total megabyte-milliseconds taken by all map tasks=31676416
        Map-Reduce Framework
                Map input records=2468572
                Map output records=2468572
                Input split bytes=87
                Spilled Records=0
                Failed Shuffles=0
                Merged Map outputs=0
                GC time elapsed (ms)=234
                CPU time spent (ms)=32140
                Physical memory (bytes) snapshot=417554432
                Virtual memory (bytes) snapshot=2801614848
                Total committed heap usage (bytes)=388497408
        File Input Format Counters
                Bytes Read=0
        File Output Format Counters
                Bytes Written=522660639
20/11/24 15:34:12 INFO mapreduce.ImportJobBase: Transferred 498.448 MB in 54.1614 seconds (9.203 MB/sec)
20/11/24 15:34:12 INFO mapreduce.ImportJobBase: Retrieved 2468572 records.
```